Question Answering

You will be implementing an encoder this week. Last week you implemented the decoder. So here it is:

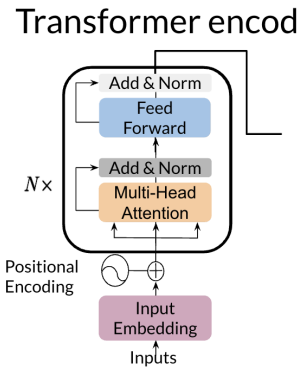

Feedforward:

```
[
      LayerNorm,
      dense,
      activation,
      dropout_middle,
      dense,
      dropout_final
]
```

Encoder block:

```
[
      Residual(
          LayerNorm,
          attention,
          dropout_,
      ),
      Residual(
          feed_forward,
      )
]
```

You can see there is a feedforward and the encoder-block above. It makes use of two residual connections, layer normalization, and dropout.

The steps you will follow to implement it are:

- Load a pre-trained model
- Process data to get the required inputs and outputs: "question: Q context: C" as input and "A" as target
- Fine tune your model on the new task and input
- Predict using your own model

Mark as completed